



SmartStock AI

An intelligent inventory management system leveraging
FastAPI, Next.js, and EfficientNetB0.



FastAPI



Next.js 14



TensorFlow



MongoDB





Backend Architecture

Framework



AI Integration



Loads model once. Handles real-time prediction from raw bytes.



Model Surgery

`routes_train.py`

Orchestrates incremental learning.

Existing APIs




Database Architecture



Strategy: MongoDB (NoSQL)

Justification: Pet products vary wildly (e.g., 'water_type' for fish vs. 'breed_size' for dogs). MongoDB's schema-less nature handles this polymorphism efficiently without sparse SQL tables.

Image Storage Process



Images are stored as **binary data** directly in the database or via GridFS. The backend provides a dedicated endpoint to serve raw bytes to the frontend, ensuring data locality.


Product Schema (JSON Preview)

```
1 {  
2   "_id": "ObjectId(...)",  
3   "name": "Premium Dog Food",  
4   "type": "Dog",  
5   "breed": "Golden Retriever",  
6   "price_predicted": 45.99,  
7   "price_modified": null,  
8   "attributes": { // Flexible schema  
9     "weight_kg": 12.5,  
10    "vaccinated": true  
11  }  
12 }
```




Classification Model: EfficientNetB0

Why EfficientNetB0?



Efficiency

Achieves high accuracy with significantly fewer parameters than ResNet50 or VGG16.




Speed

Faster inference time is critical for the real-time web application experience.

Transfer Learning


Leverages features learned from ImageNet to train effectively on smaller pet datasets.



ImageNet Features

Smaller Pet Datasets

Architecture



Training Data Strategy

Data Directory Structure



Images are organized hierarchically: Category > Breed.
The metadata used for the price regression model is stored separately in CSV format.

Augmentation



Techniques: Rotation, Zoom, Horizontal Flip.
Increases dataset diversity and prevents overfitting by ensuring the model learns features (fur, ears) rather than specific lighting or poses.

Hyperparameters



Batch Size: 32 000

Phase 1: 50 Epochs (Head Only)

50

Phase 2: 10 Epochs (Fine-Tune)

10



Dynamic Training Strategies

Model Surgery

Goal: Add a new breed (e.g., Poodle) without retraining from scratch.

1. Prune & Expand:

Remove output layer (N neurons). Create new layer with N+1 neurons.



2. Weight Transfer:

Copy old weights. Initialize only the new neuron.

Fine-Tuning

Goal: Improve accuracy for an existing breed using new photos.



Low Learning Rate:

Uses 1e-5 to gently adjust weights without destroying pre-learned features.



Rehearsal Strategy:

Mixes in data from other breeds to prevent "Catastrophic Forgetting" of old knowledge.

Frontend Engineering

Framework & Stack




Key UI Features



⚖️ Model Selection: Price Regression

1. Chosen Technique & 2. Expected Outcome



3. Justification: Why Neural Network?

Why use a Neural Network over XGBoost or Random Forest?

Criterion	Random Forest	Neural Network (Chosen)
Non-Linearity	Good/Trees	Excellent/Activations
Ecosystem	Scikit-Learn	TensorFlow(Unified)
Complexity	Lower	High Capacity





Price Model: Synthetic Data Engine



1. The Logic Engine

Since real-world pet price data is unstructured, we use a Python script (`generate_price_dataset.py`) to apply market rules to existing images.



Base Characteristics

Lookups per breed: Base Price (e.g., \$1500) and Avg Weight (e.g., 30kg).



Price Formula

Final Price = Base × (Age Factor) + Health Bonus + Origin Premium + Random Noise



2. Generated Output Schema

The script outputs a CSV file with 7 features per sample:

Type / Breed	Age	Weight	Health	Origin	Target Variable:
"Dog", "Golden Retriever"	12 Months (Younger = \$\$\$)	29.5 kg	2 (Excellent)	"Scotland" (Premium)	Price: ↑ \$1,850.50

Preparing Data for AI

1. Feature Engineering



A. Normalization (Scaling) Numerical

Ensures large numbers don't dominate the model.

Age (30 mo)

$\div 60 \text{ (Max)}$

0.50

Weight (10 kg)

$\div 50 \text{ (Max)}$

0.20



B. Encoding (LabelEncoder) Categorical

Converts text labels into unique ID numbers.

Breed:
"Poodle"

Lookup
ID

28

Country:
"USA"

Lookup
ID

1

2. Network Architecture

Input Layer (7
Neurons)

Deep Hidden Layers

Dense(128) + BatchNorm + Dropout

Dense(64) + BatchNorm + Dropout

Dense(32)

Output Layer (1
Neuron)

Linear Activation = Predicted Price (\$)

?

Technical Q&A: Core Decisions

Frontend Architecture

Why Next.js (SSR) instead of standard React?

We chose **Server-Side Rendering** to improve SEO and initial load times. For an e-commerce platform, ensuring product pages are indexable by search engines is critical, which a client-side SPA (Single Page App) often struggles with.



Price Regression

Why separate the Price Model from the Image?

Signal-to-Noise Ratio. A photo doesn't show "Vaccination Status" or "Pedigree." A dedicated MLP using structured metadata (Age, Health, Origin) yields far higher regression accuracy than attempting to estimate value from pixels alone.